




Universal Support for Scoped Memory Access Instrumentation

Chris Kjellqvist, Pantea Zardoshti, Michael Spear

- » Advances in software and hardware research brings new memory features
 - » Transactional Memory
 - » Non-Volatile Memory (NVM)
 - » Non Uniform Memory Access (NUMA)
 - » Guarded Memory



Intel Coffee Lake Die Layout

- » A transaction is a sequence of steps executed by a single thread.
- » Atomic transaction
 - » Commit: memory writes take effect
 - » Abort: effects rolled back

```
void f(int* n){  
    if (*n == 4)  
        *n = *n + 1;  
}  
int a = 4;  
f (&a);
```

» Consider the previous function, but when two threads execute the function at the same time

*n	Thread 1	Thread 2
4		if (*n == 4)
4	if (*n == 4)	
5	*n = *n + 1;	
6		*n = *n + 1;

```
void f(int* n){  
    if (*n == 4)  
        *n = *n + 1;  
}  
//initialized globally as 4  
f (some_shared_ptr);
```

» We can make our code work by using transactions

Thread 1
`transaction {`

`if (*n == 4)`

`*n = *n + 1;`

`} (transaction`
`aborts)`

Thread 2

`transaction {`

`if (*n == 4)`

`*n = *n + 1;`

`} (transaction`
`succeeds)`

```
void f(int* n){  
    transaction {  
        if (*n == 4)  
            *n = *n + 1;  
    }  
}  
//initialized globally as 4  
f (some_shared_ptr);
```

» However using current transactional memory libraries is a much more complicated process than previously shown

» Static Separation

RSTM API

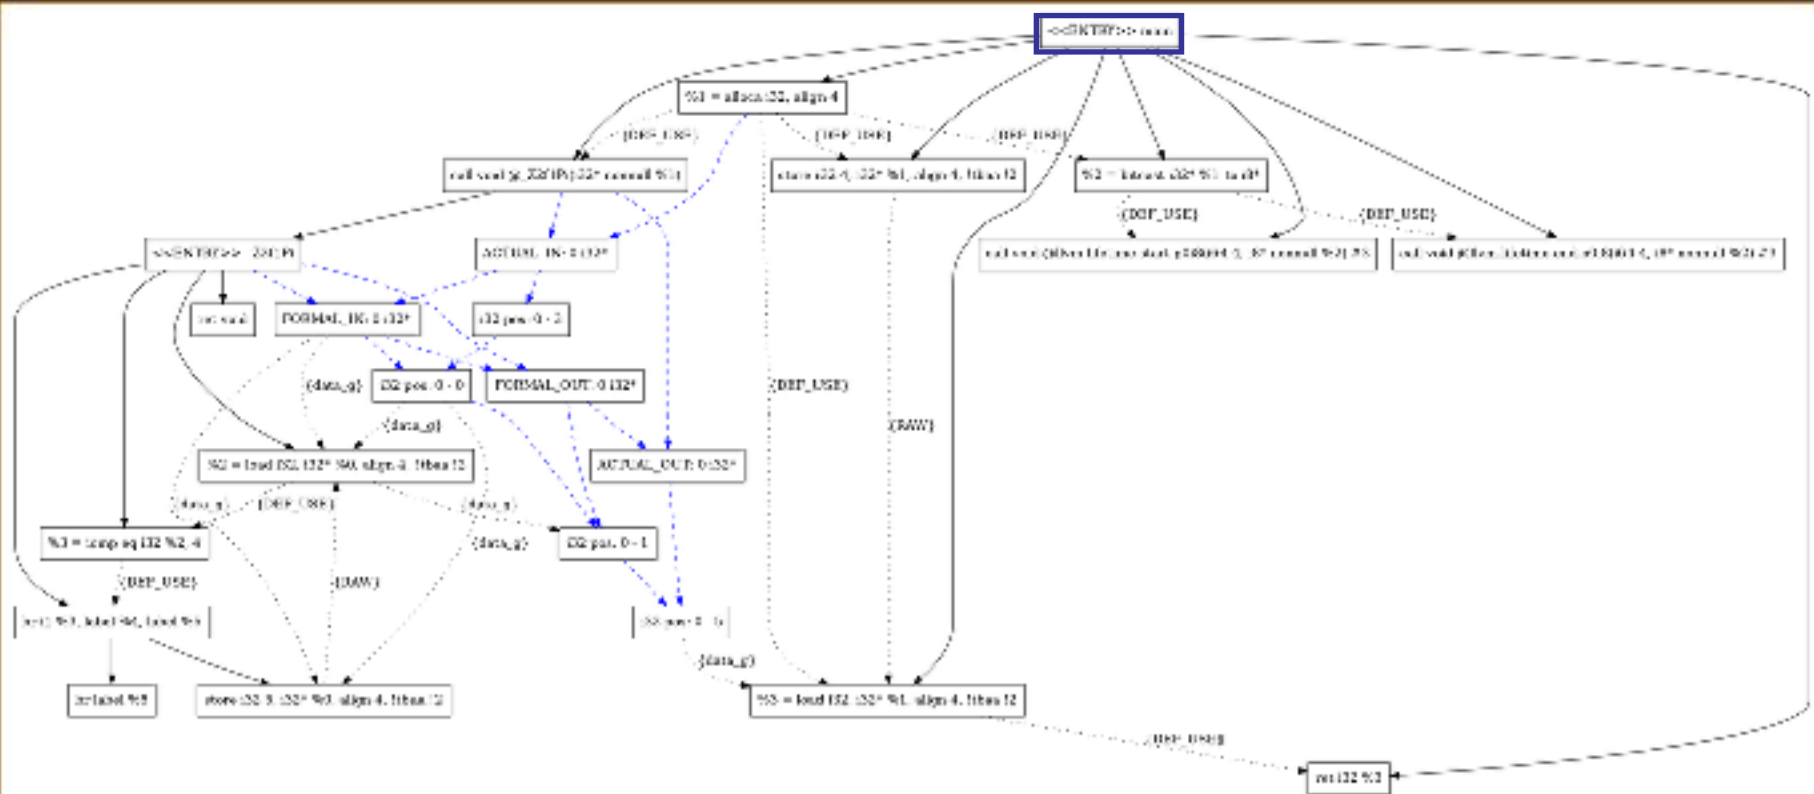
```
void f(int* n){  
    TM_BEGIN(atomic) {  
        int z = TM_READ(n);  
        if (z == 4)  
            TM_WRITE(n, z+1);  
    } TM_END;  
}
```

Intel TSX API

```
void f(int* n){  
    int status;  
    while ((status = _xbegin()) != _XBEGIN_STARTED) {  
        if (*n == 4)  
            *n = *n + 1;  
        _xend()  
    }  
}
```

- » To accomplish this task, we have developed a set of C++ keywords to streamline the use of TM
 - » Annotate functions that you'd like to act as transactions
 - » Annotate variables in a similar way

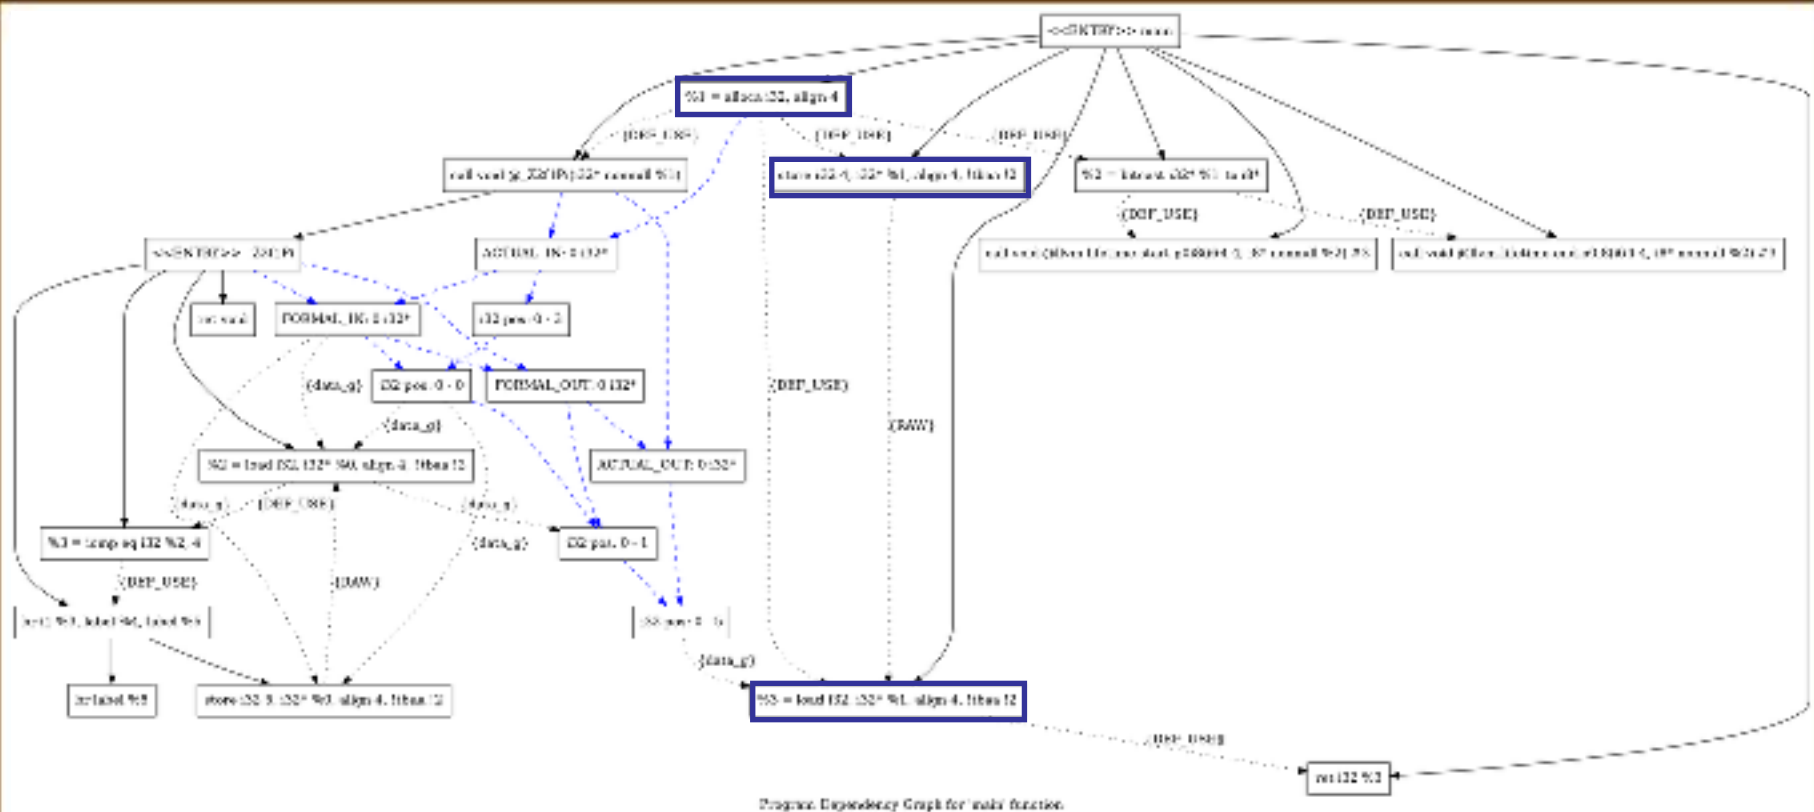
```
TX_SAFE void f1(int* n){  
    if (*n == 4)  
        *n = *n + 1;  
}  
int main() {  
    TX_VAR(int) a = 4;  
    TX_PTR(int) *b = malloc(4);  
  
    f1 (&a);  
    return a;  
}
```



```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    int a = 4;
    f1 (&a);
    return a;
}
```

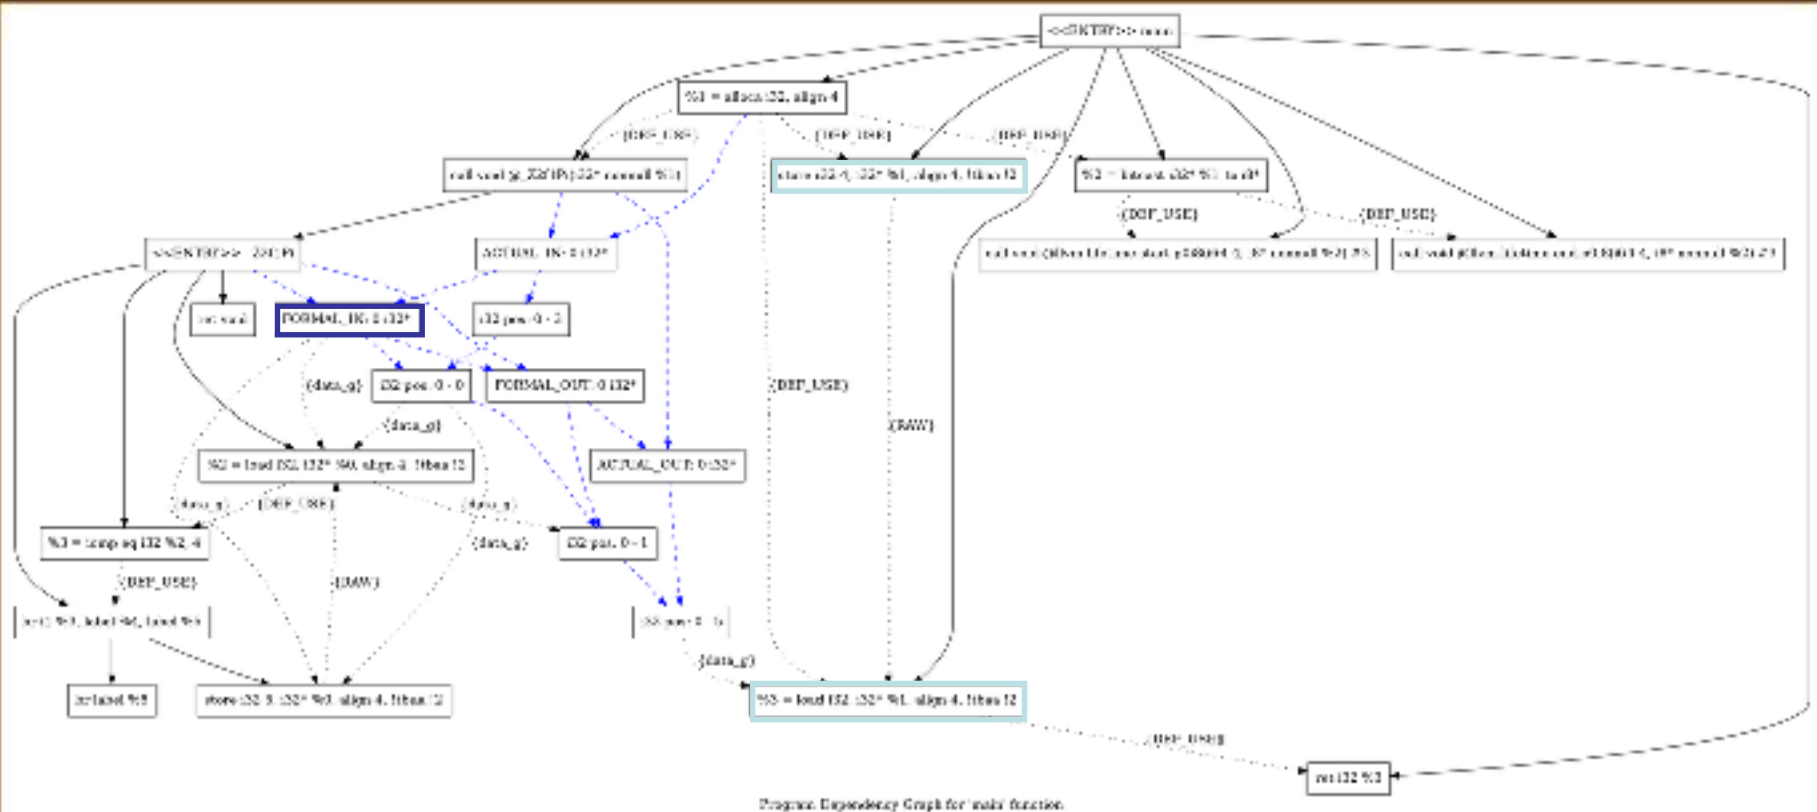




```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    int a = 4;
    f1 (&a);
    return a;
}
```

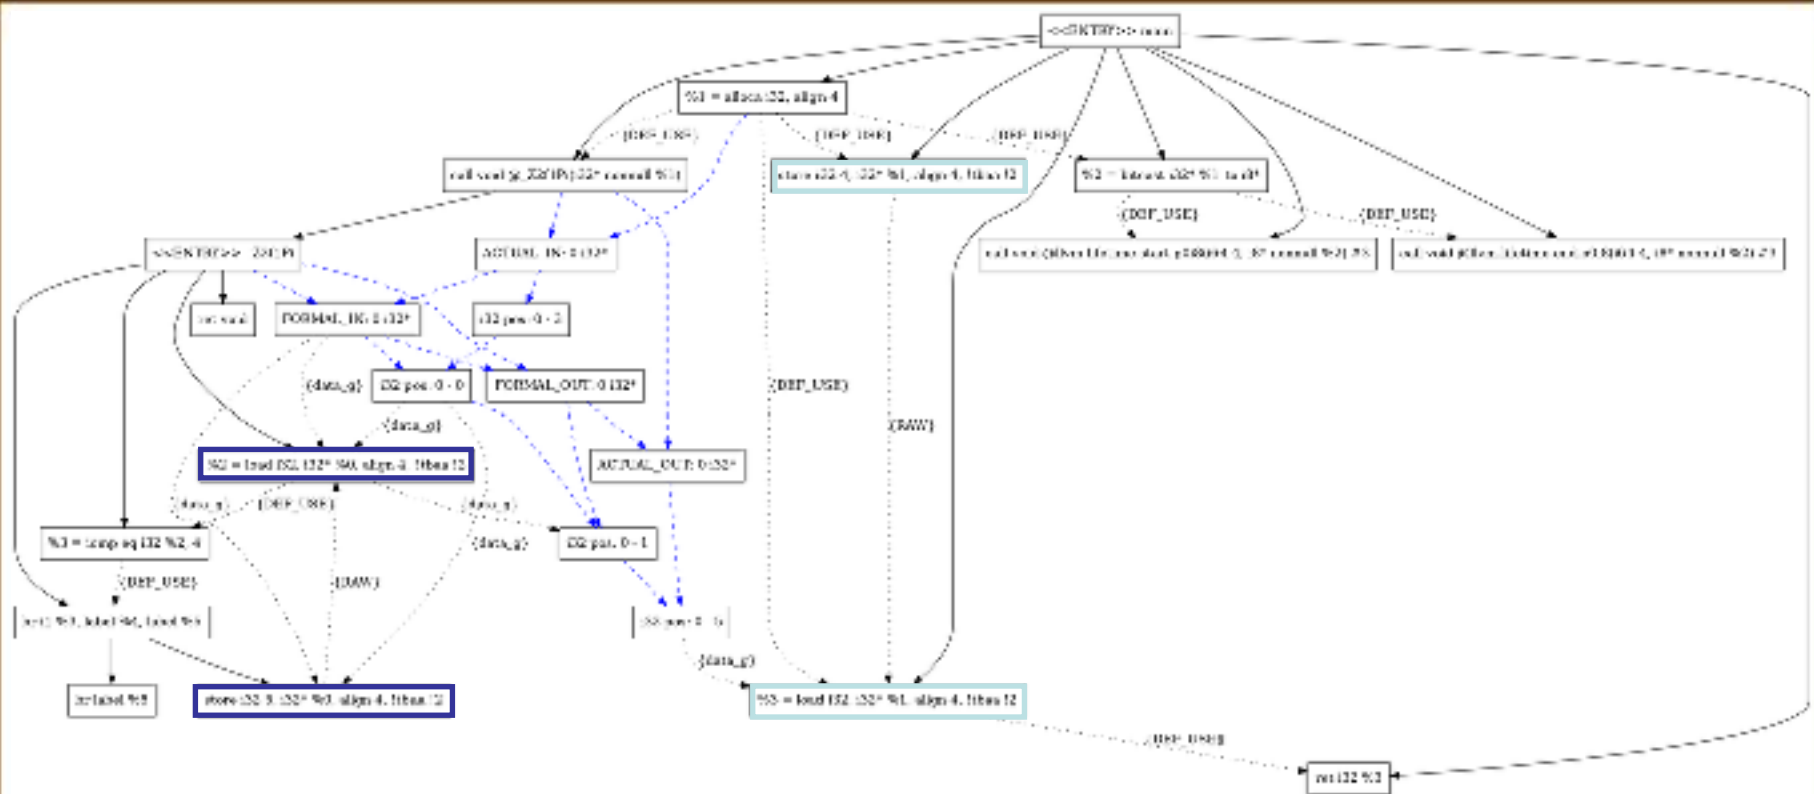




```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    int a = 4;
    f1 (&a);
    return a;
}
```

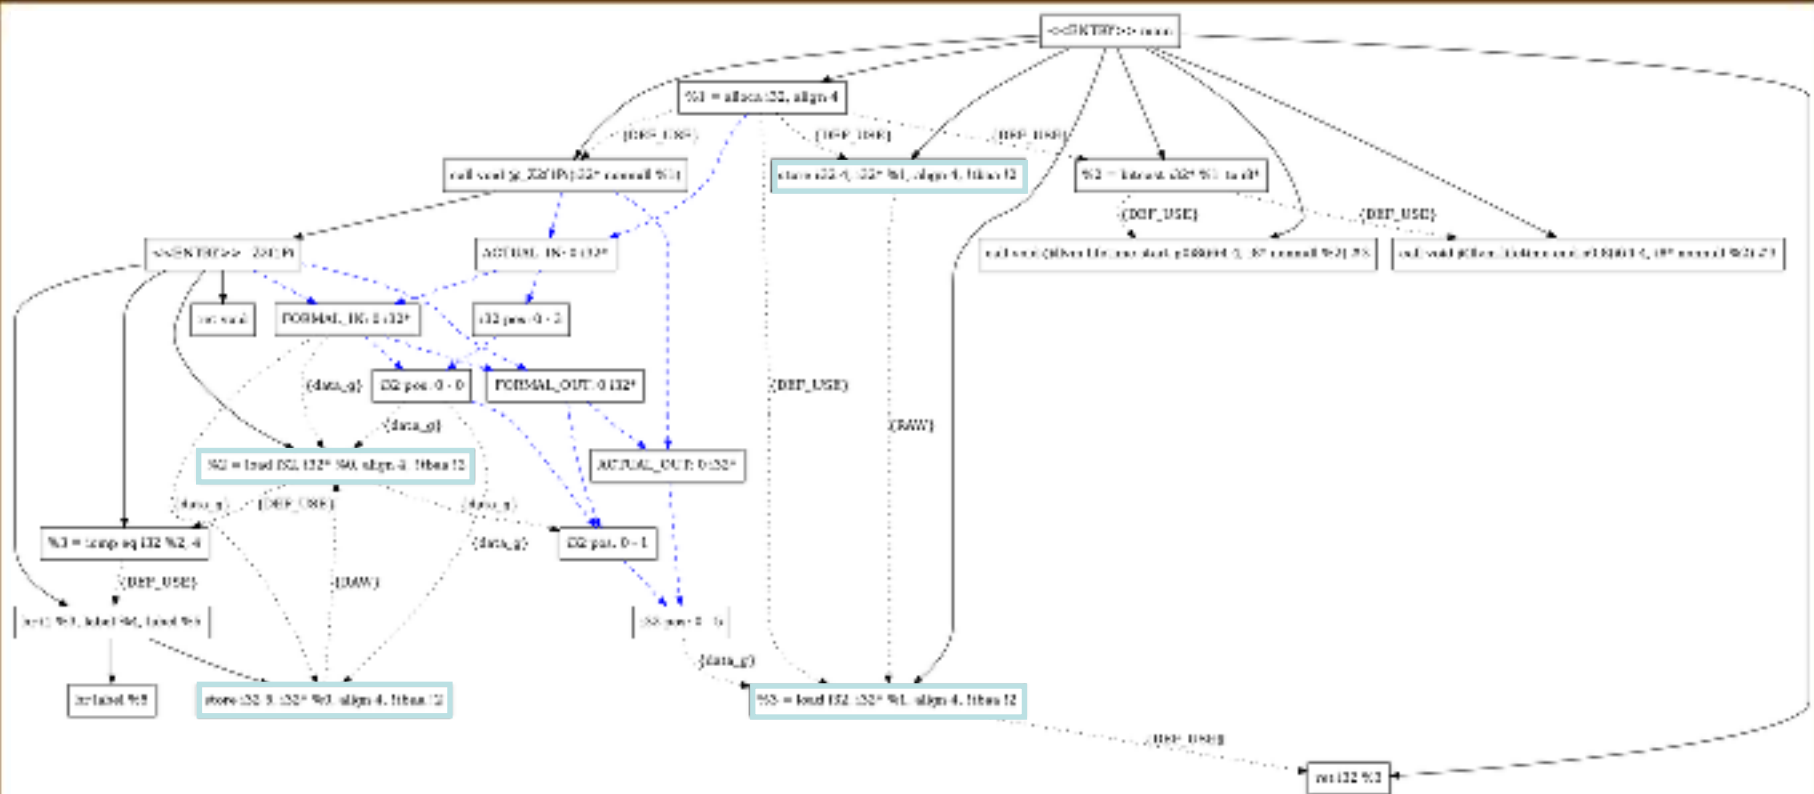




```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    int a = 4;
    f1 (&a);
    return a;
}
```

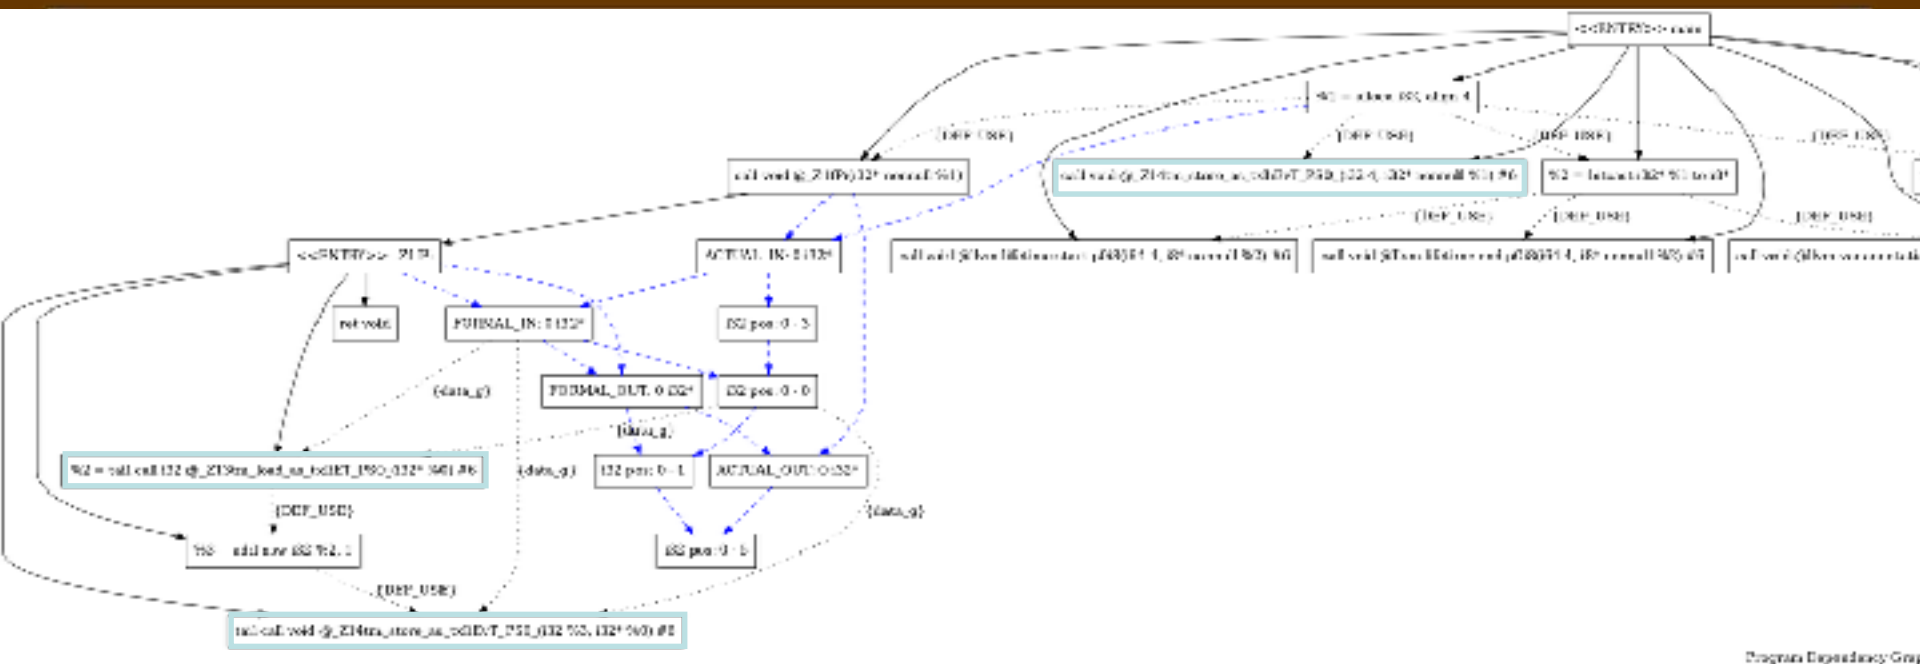




```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    int a = 4;
    f1 (&a);
    return a;
}
```





Program Dependencies Graph

```
void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    TX_VAR(int) a = 4;
    f1 (&a);
    return a;
}
```



```
int main() {
    TX_VAR(int) a = 4;
    f1 (&a);
    return a;
}
```

- » Ease of use to usage of TM libraries
- » Optimization by data region instead of code region

RSTM API

```
void f(int* n){
    TM_BEGIN(atomic) {
        int z = TM_READ(n);
        if (z == 4)
            TM_WRITE(n, z+1);
    } TM_END;
}
```

Intel TSX API

```
void f(int* n){
    int status;
    while ((status = _xbegin()) != _XBEGIN_STARTED) {
        if (*n == 4)
            *n = *n + 1;
        _xend()
    }
}
```

Before

After

```
TX_SAFE void f1(int* n){
    if (*n == 4)
        *n = *n + 1;
}
```

```
int main() {
    TX_VAR(int) a = 4;
    f1 (&a);
    return a;
}
```


- » Transactions can reduce assumptions of multithreaded code
- » Heterogenous memory systems are hard to use consistently
- » Variable attributes let us transform loads/stores outside of transactions to assist with static separation